

Principles of Robot Autonomy I

PnP, SfM, RANSAC




Logistics

- Homework 3: due Tues, Oct 28
- Homework 4: out Tues Nov 4 (**After the Midterm**)
- Midterm window: Wed, Oct 29, 5pm – **Sat, Nov 1**
 - Take home, 72 hour window
 - Check out exam on gradescope
 - You will have personal 5 hour time slot
 - Open notes, book, HW solutions
 - No internet, no GenAI, no working with others
- Lecture 9:
 - N point perspective problem (PnP)
 - Structure from Motion (SfM)
 - RANSAC and feature correspondences

Robot Perception

See: Perception Stack




Perception Stack: Computer vision, filtering, SLAM

Use sensor data to update the models


Localization, Mapping, Tracking:

- EKF/Monte Carlo localization
- Occupancy grid mapping
- Pose graph optimization
- Tracking (EKF and Particle Filter)
- AA273: Filtering (Schwager)
- AA275: Navigation (Gao)



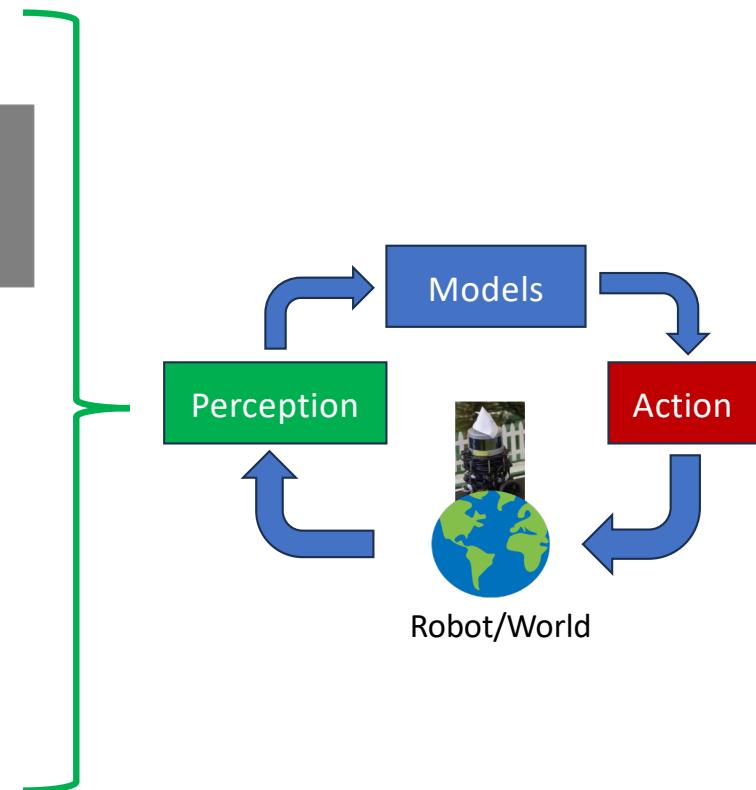
Information extraction

- Computer vision: features, correspondences, Structure from Motion (SfM), depth
- Lidar scan matching, ICP
- CS231A: Comp Vision



Sensors:


- RGB Cameras, RGB-D/stereo cameras, Lidar
- IMU, GPS, wheel encoders



Recall: Camera calibration

Goal: find the intrinsic and extrinsic parameters of the camera

- Known 3D world and 2D images coordinates
- Known correspondences between them



Credit: FP Chapter 1

Solution: take SVD and find singular vector with smallest singular value.

$$p^h = K[R \quad t]P_W^h$$




Projection matrix M

From world frame to camera frame (inverse cam pose)

$$\min_{m \in R^{12}} \|\tilde{P}m\|^2$$

$$\text{subject to } \|m\|^2 = 1$$

Examples for Calibration Images



- Due to co-planarity of cal board, need multiple images from different poses
- Origin of world frame defined by board lower right corner

Source: Wikipedia

Perspective-n-Point (PnP) Problem (e.g. for object pose tracking)

Goal: find 3D object pose with respect to camera


- Known 3D coords of object points in its body frame, 2D image coordinates, and known correspondences
- Known camera calibration matrix K
- **Same solution technique as camera calibration!**

Pose of object wrt camera

$$K^{-1}p^h = [R \quad t]P_W^h$$


$$\begin{aligned} & \min_{m \in R^{12}} \|\tilde{P}m\|^2 \\ & \text{subject to } \|m\|^2 = 1 \end{aligned}$$

- \tilde{P} and m defined with K^{-1} on LHS.



Solution: take SVD and find singular vector with smallest singular value.

Structure from motion (SFM)




Given m images of n fixed 3D points

$$p_{j,k}^h = M_k P_j^h$$

Find:

- m projection matrices M_k (**motion**)
- n 3D points P_j (**structure**)

Measuring depth



$$p^h = K[R \ t]P_W^h$$

Homogeneous coordinates

- Once the camera is calibrated, can we measure the location of a point P in 3D given its known observation p ?
- No: one can only say that P is located *somewhere* along the line joining p and O !


Issues with recovering structure (3D point locations in world frame)



Recovering structure


- **Structure:** 3D point locations to be reconstructed by having access to only 2D images
- Common methods
 1. Structure from motion (SfM): processes two images taken with the same or different cameras at *different times* and from different *unknown* positions
 2. Stereo vision: processes two distinct images taken at the *same time* and assumes that the relative pose between the two cameras is *known*
 3. Monocular depth estimation: Deep learned from context using a history of depth images (learned scale of objects, lighting queues, function queues, ect)
 4. Depth from focus: determines distance to one point by taking multiple images with better and better focus

Epipolar geometry



- Consider image coords p and p' of a point P in world frame observed by two cameras centered at O, O'
- These five points all belong to the *epipolar plane* defined by p, O, O' , or equivalently, p', O, O'
- **Epipolar constraint:** potential matches for p must lie on epipolar line l' (and vice-versa)

Epipolar constraint derivation



- Epipolar constraint: $\overline{O'p'}$, \overline{Op} , and $\overline{O'O}$ (when expressed in that same frame) must be co-planar, or

$$\overline{O'p'} \cdot [\overline{O'O} \times \overline{Op}] = 0$$

Aside: matrix notation for cross product

- Cross product can be expressed as the product of a **skew-symmetric** matrix and a vector


$$a \times b = \underbrace{\begin{bmatrix} 0 & -a_3 & a_2 \\ a_3 & 0 & -a_1 \\ -a_2 & a_1 & 0 \end{bmatrix}}_{:= [a]_{\times}} \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix} = [a]_{\times} b$$

Epipolar constraint derivation

(more details: CS231a)

- Translate into vector expressions, all in frame O'

$$\overline{O'p'} \cdot [\overline{O'O} \times \overline{Op}] = 0$$
$$(K'^{-1}p')^T [t]_{\times} (RK^{-1}p) = 0$$




$$p'^T F p = 0 \quad \text{where} \quad F = K'^{-T} [t]_{\times} R K^{-1}$$

Fundamental matrix

Key facts

- F is referred to as the **fundamental matrix**
- $l' = Fp$ (resp. $l = F^T p'$) represents the epipolar line corresponding to the point p (resp. p') in the second (resp. first) image. This exploits the homogenous notation for lines.
- $Fe = F^T e' = 0 \rightarrow F$ is also singular (as t is parallel to the coordinate vectors of the epipoles)
- F has 7 DoF (9 elements – common scaling – $\det(F)=0$)

Usefulness of fundamental matrix



- Assume F is given
- Given a point in image 1, one can compute the corresponding epipolar line in image 2 **without any additional information needed!**

Estimating the fundamental matrix

- 8-point algorithm

$$\begin{aligned}
 p = \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \quad p' = \begin{bmatrix} u' \\ v' \\ 1 \end{bmatrix} & \Rightarrow \begin{bmatrix} u' & v' & 1 \end{bmatrix} \begin{bmatrix} F_{11} & F_{12} & F_{13} \\ F_{21} & F_{22} & F_{23} \\ F_{31} & F_{32} & F_{33} \end{bmatrix} \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = 0 \\
 \Rightarrow \begin{bmatrix} u'u & v'u & u & u'v & v'v & v & u' & v' & 1 \end{bmatrix} \begin{bmatrix} F_{11} \\ F_{21} \\ F_{31} \\ F_{12} \\ F_{22} \\ F_{32} \\ F_{13} \\ F_{23} \\ F_{33} \end{bmatrix} = 0 & \Rightarrow Wf = 0
 \end{aligned}$$

nx9 matrix of known coefficients
 f

$$\min_{f \in R^9} \|Wf\|^2 \Rightarrow \tilde{F}$$

subject to $\|f\|^2 = 1$

- Given $n \geq 8$ correspondences, one then solves

Enforcing the rank constraint

- \tilde{F} satisfies the epipolar constraints, but is not necessarily singular (hence, is not necessarily a proper fundamental matrix)
- Enforce rank constraint (again, via SVD decomposition)

$$\begin{aligned} \text{Find } F \text{ that minimizes } \|F - \tilde{F}\|^2 &\quad \leftarrow \text{Frobenius norm} \\ \text{subject to } \det(F) = 0 \end{aligned}$$

- 8-point algorithm
 1. Use least norms to compute \tilde{F} by taking SVD of W , and finding right singular vector with smallest singular value.
 2. Enforce rank constraint on \tilde{F} by taking SVD and setting smallest two singular values to zero, then re-multiply $\tilde{F} = U\Sigma V^T$

Obtaining relative camera pose

- Recall:

$$F = K'^{-T} [t]_{\times} R K^{-1}$$


- To get (R, t) we compute

$$[t]_{\times} R = K'^T F K$$

- Use RQ decomposition to get

$$[t]_{\times} R \rightarrow ([t]_{\times}, R)$$

$$[t]_{\times} \rightarrow t$$




Solution to SFM problem (high-level)


- Several approaches available:
 - Algebraic approach (by fundamental matrix)
 - Bundle adjustment (large-scale non-convex optimization problem to find $[R \ t]$ camera poses for each image, and 3D points (i.e., structure) in common world frame)
- Algebraic approach (2-views) – gives relative poses between two successive views, and 3D points relative to first view
 1. Compute fundamental matrix F (e.g., via 8-point algorithm)
 2. Use F to estimate camera projection matrices
 3. Use camera projection matrices for triangulation
- Used as "front end" in visual SLAM, to obtain pose graph for PGO back end

Application of SFM: visual odometry


- **SLAM/Visual Odometry**: estimate the motion of the robot by using visual input (and possibly additional information)
 - Single camera: absolute scale must be estimated in other ways
 - Stereo camera: measurements are directly provided in absolute scale



Elephant in the room: how to find correspondences for PnP, stereo, or SfM?




Features, feature descriptors, feature matching, RANSAC



Keypoint Features and their Descriptors

- **Goal:** *describe* keypoints so that we can compare them across images or use them for object detection or matching
- Desired properties:
 - Invariance with respect to pose, scale, illumination, etc.
 - Distinctiveness



RANSAC iterations

- In principle, one would need to check all possible combinations of 2 points in dataset
- If $|S| = N$, number of combinations is $\frac{N(N-1)}{2} \rightarrow$ too many
- However, if we have a rough estimate of the percentage of inliers, we do not need to check all combinations...

E.g., RANSAC for line fitting (for simplicity)

Data: Set S consisting of all N points


Result: Set with maximum number of inliers
(and corresponding fitting line)

while $i \leq k$ **do**

randomly select 2 points from S ;
fit line l_i through the 2 points;
compute distance of all other points to line l_i ;
construct *inlier* set, i.e., count number of
points with distance to the line less than γ ;
store line l_i and associated set of inliers;
 $i \leftarrow i + 1$

end

Choose set with maximum number of inliers



RANSAC

Data: Set S consisting of all N points


Result: Set with maximum number of inliers
(and corresponding fitting line)

while $i \leq k$ **do**

 → randomly select 2 points from S ;
 fit line l_i through the 2 points;
 compute distance of all other points to line l_i ;
 construct *inlier* set, i.e., count number of
 points with distance to the line less than γ ;
 store line l_i and associated set of inliers;
 $i \leftarrow i + 1$

end

Choose set with maximum number of inliers



RANSAC

Data: Set S consisting of all N points


Result: Set with maximum number of inliers
(and corresponding fitting line)

while $i \leq k$ **do**

 randomly select 2 points from S ;
 fit line l_i through the 2 points;
 compute distance of all other points to line l_i ;
 construct *inlier* set, i.e., count number of
 points with distance to the line less than γ ;
 store line l_i and associated set of inliers;
 $i \leftarrow i + 1$

end

Choose set with maximum number of inliers



RANSAC

Data: Set S consisting of all N points


Result: Set with maximum number of inliers
(and corresponding fitting line)

while $i \leq k$ **do**

randomly select 2 points from S ;
fit line l_i through the 2 points;
compute distance of all other points to line l_i ;
construct *inlier* set, i.e., count number of
points with distance to the line less than γ ;
store line l_i and associated set of inliers;
 $i \leftarrow i + 1$

end

Choose set with maximum number of inliers



RANSAC

Data: Set S consisting of all N points

Result: Set with maximum number of inliers
(and corresponding fitting line)

while $i \leq k$ **do**

 randomly select 2 points from S ;

 fit line l_i through the 2 points;

 compute distance of all other points to line l_i ;


 → construct *inlier* set, i.e., count number of
 points with distance to the line less than γ ;

 store line l_i and associated set of inliers;

$i \leftarrow i + 1$

end

Choose set with maximum number of inliers



RANSAC

Data: Set S consisting of all N points


Result: Set with maximum number of inliers
(and corresponding fitting line)

while $i \leq k$ **do**

randomly select 2 points from S ;
fit line l_i through the 2 points;
compute distance of all other points to line l_i ;
construct *inlier* set, i.e., count number of
points with distance to the line less than γ ;
store line l_i and associated set of inliers;
 $i \leftarrow i + 1$

end

Choose set with maximum number of inliers



RANSAC

Data: Set S consisting of all N points


Result: Set with maximum number of inliers
(and corresponding fitting line)

while $i \leq k$ **do**

 randomly select 2 points from S ;
 fit line l_i through the 2 points;
 compute distance of all other points to line l_i ;
 construct *inlier* set, i.e., count number of
 points with distance to the line less than γ ;
 store line l_i and associated set of inliers;
 $i \leftarrow i + 1$

end

Choose set with maximum number of inliers



RANSAC

Data: Set S consisting of all N points

Result: Set with maximum number of inliers
(and corresponding fitting line)

while $i \leq k$ **do**


 randomly select 2 points from S ;
 fit line l_i through the 2 points;
 compute distance of all other points to line l_i ;
 construct *inlier* set, i.e., count number of
 points with distance to the line less than γ ;

 store line l_i and associated set of inliers;

$i \leftarrow i + 1$

end

Choose set with maximum number of inliers



Core “feature based” CV pipeline

- Input two images (stereo or SfM), or image and 3D model (PnP)
- Feature detection, feature matching -> N correspondences
- **Problem: in practice, we get loads of bad feature matches**
- Use Random Sample Consensus (RANSAC) to reject outliers
- In RANSAC, we
 - Sample a small subset n of matched features
 - **Solve geometric CV problem with these correspondences (PnP, stereo, SfM)**
 - Check solution fit with remaining (N-n) feature matches – record fit score
 - Repeat
 - Return the fit with the highest score