

Principles of Robot Autonomy I

Occupancy grid mapping , Frontier Exploration,
Bayesian estimation and filtering,

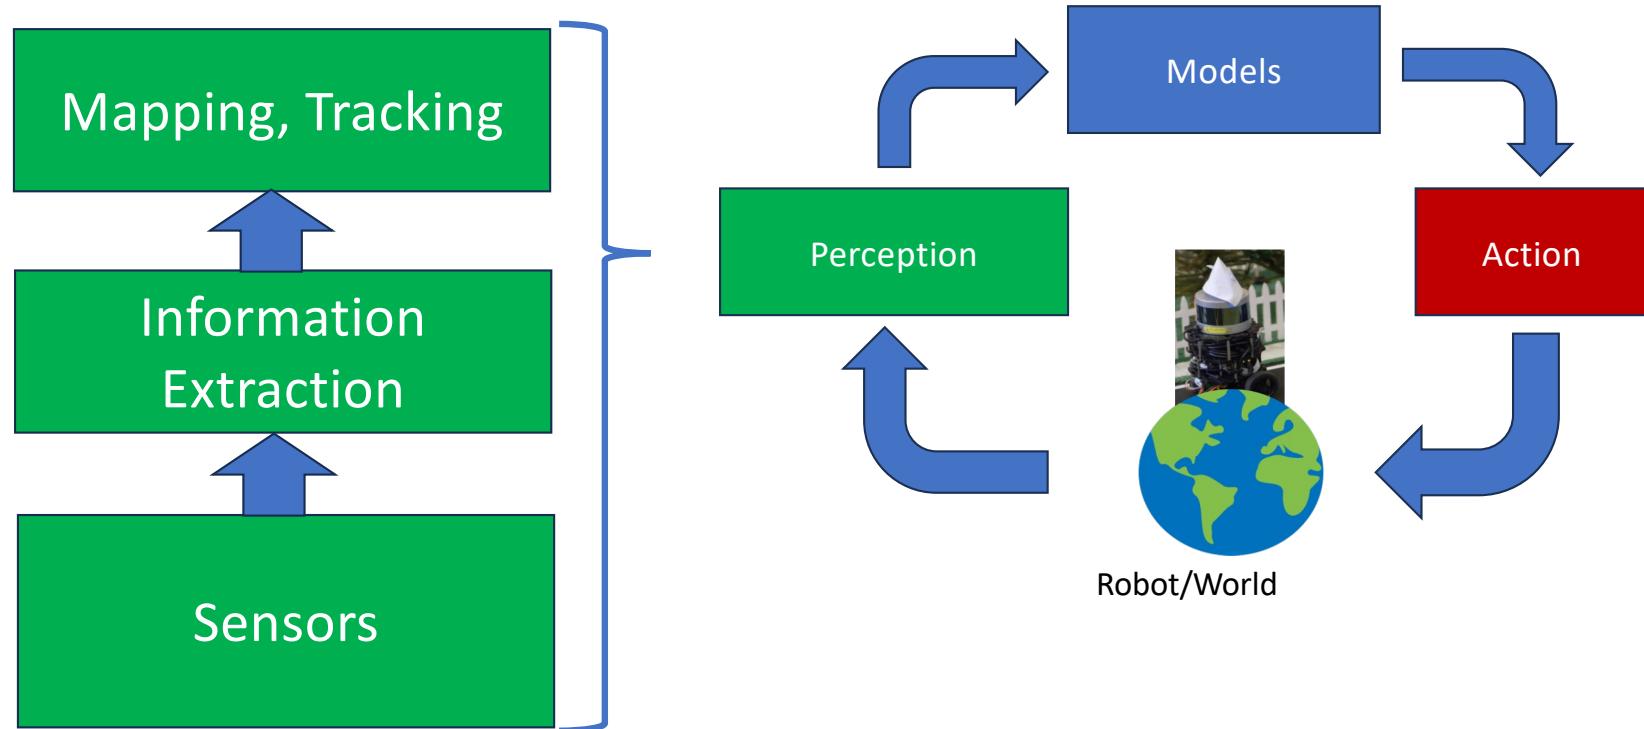


Logistics

- Homework 4: due Thurs 11/13
- Homework 5: out Thurs 11/13, Due 12/2 (last one!)
- Guest Lecture: Dr. Vincent Vincent Vanhoucke, Distinguished Engineer, Waymo, Thurs 11/20
- Midterm grades out soon
- Lecture 14:
 - Occupancy grid mapping
 - Frontier exploration, information gathering
 - General Bayesian filter

Robot Perception

Perception Stack

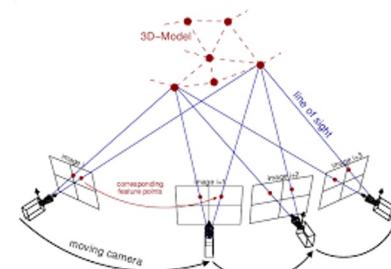
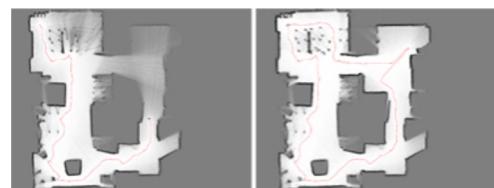


Perception Stack: Computer vision, filtering, SLAM

Use sensor data to update the models

Localization, Mapping, Tracking:

- EKF/Monte Carlo localization
- Occupancy grid mapping
- Factor graphs/SLAM
- Tracking (EKF and Particle Filter)
- AA273: Filtering (Schwager)
- AA275: Navigation (Gao)



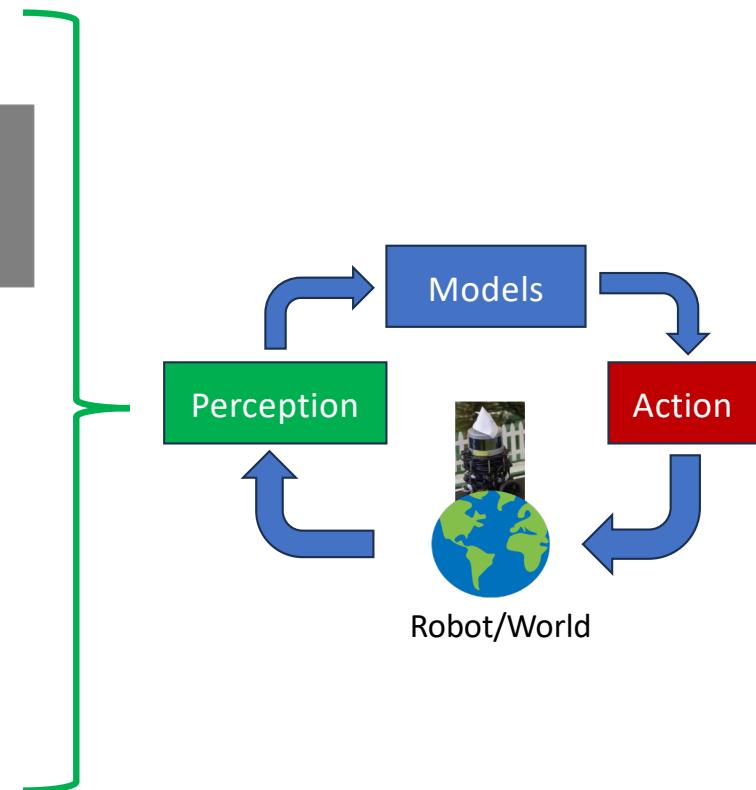
Information extraction

- Computer vision: features, correspondences, Structure from Motion (SfM), depth
- Lidar scan matching, ICP
- CS231A: Comp Vision



Sensors:

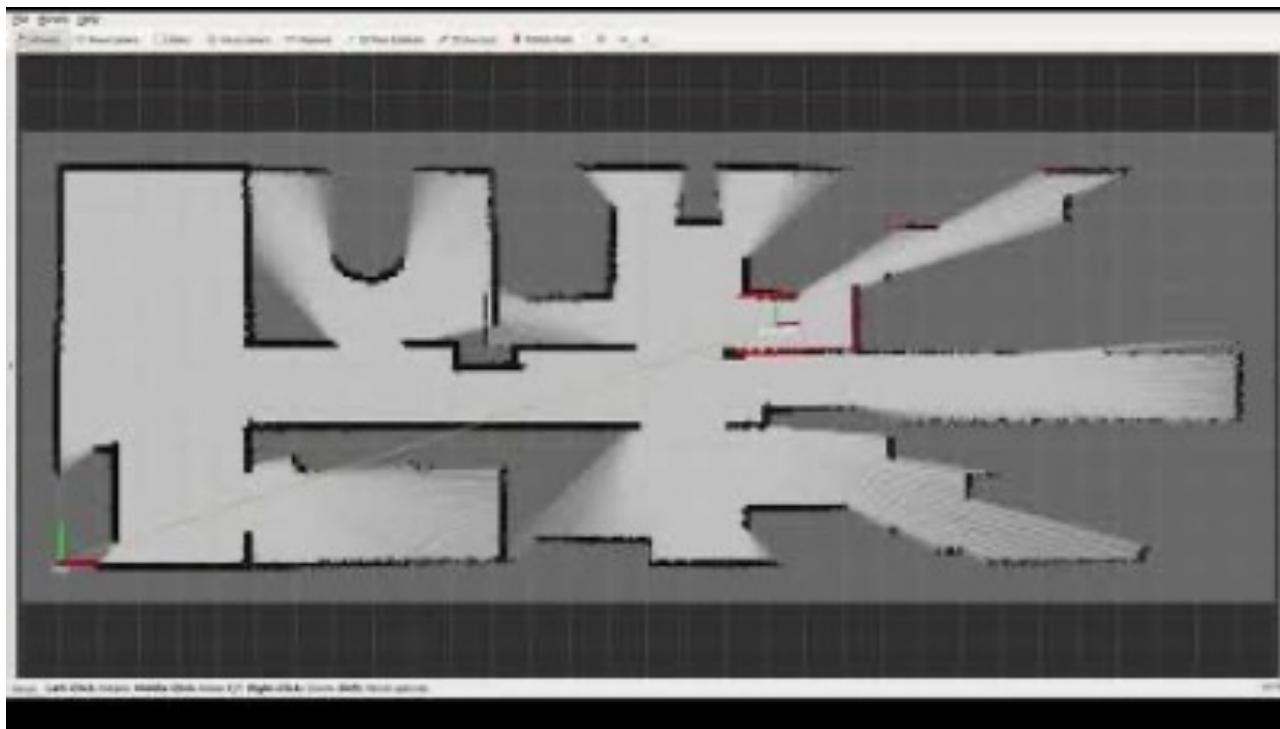
- RGB Cameras, RGB-D/stereo cameras, Lidar
- IMU, GPS, wheel encoders



Today's lecture

- Aim
 - Learn basic concepts about Bayesian filtering and apply to occupancy grid mapping and Frontier Exploration.
- Readings
 - S. Thrun, W. Burgard, and D. Fox. Probabilistic robotics. MIT press, 2005.
Chapter 2, Chapter 9

Example: Occupancy grid mapping



Probabilistic occupancy grid mapping

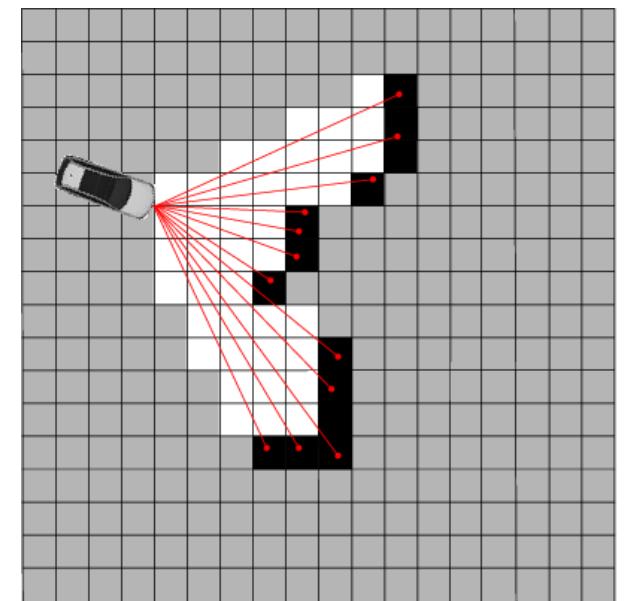
Occupancy map: $m \in \{0, 1\}^{M \times N}$

Single grid cell: $m_{i,j} \in \{0, 1\}$

Measurement model (Lidar, stereo, vision):

$$p(z_t(i,j) | m(i,j)) = \begin{bmatrix} \text{True negative} & \text{False negative} \\ p(0 | 0) & p(0 | 1) \\ p(1 | 0) & p(1 | 1) \\ \text{False positive} & \text{True positive} \end{bmatrix}$$

See point in cell Cell occupied



Cell (i,j) only gets a measurement if it is inside the footprint of sensor

Footprint depends on robot pose (assumed known)

Probabilistic occupancy grid mapping

E.g.

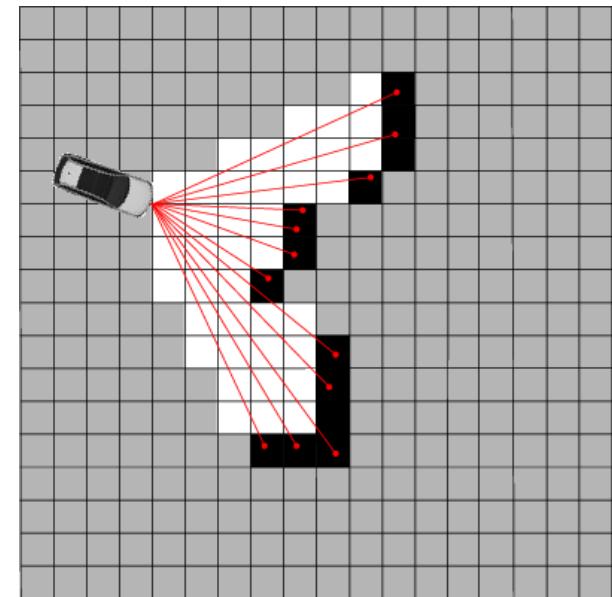
$$p(z(i, j) \mid m(i, j)) = \begin{bmatrix} \text{True negative} & \text{False negative} \\ 0.9 & 0.05 \\ 0.1 & 0.95 \\ \text{False positive} & \text{True positive} \end{bmatrix}$$

Goal:

Find probability distribution over each occupancy cell given sequence of measurements

$$p(m(i, j) \mid z_{1:t}(i, j))$$

Sequence of measurements from 1 to t



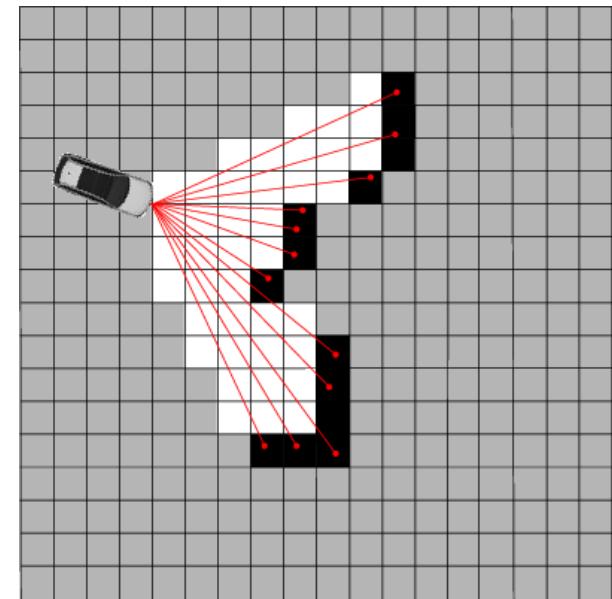
Probabilistic occupancy grid mapping

Recall:

$$p(m(i, j) = 1 \mid z_{1:t}(i, j)) \in [0, 1]$$

Called a binary random variable,
or a Bernoulli random variable

Each cell in the map is a *probability*
Ground truth map is binary, but unknown.



Bayes' rule and probabilistic inference

- Assume x is a quantity we would like to infer from y
- Bayes rule allows us to do so through the inverse probability, which specifies the probability of data y assuming that x was the cause

Posterior probability distribution



$$p(x | y) = \frac{p(y | x)p(x)}{\int p(y | x')p(x') dx'}$$

Prior probability distribution

Normalizer, does not depend on $x := \eta^{-1}$

- Notational simplification

$$p(x | y) = \eta p(y | x)p(x)$$

Independence vs Conditional Independence

- The following are equivalent: X and Y are independent

$$p(x, y) = p(x)p(y)$$

$$p(x | y) = p(x)$$

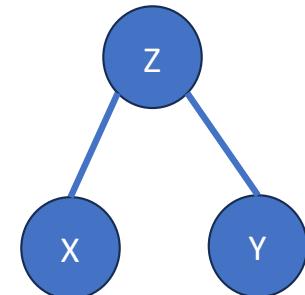
$$p(y | x) = p(y)$$

- X and Y are conditionally independent given Z

$$p(x, y | z) = p(x | z)p(y | z)$$

$$p(x | y, z) = p(x | z)$$

$$p(y | x, z) = p(y | z)$$



- “Puddles” and “umbrellas” given “raining”

Sequential Bayes' with Conditionally Independent Measurements

Sequence of measurements: $z_{1:t}$

Unknown map state: x

Conditionally Indep. measurements: $p(z_{1:t} \mid x) = \prod_{\tau=1}^t p(z_\tau \mid x)$

Bayes' Rule: $p(x \mid z_{1:t}) = \eta_t \prod_{\tau=1}^t p(z_\tau \mid x)p(x)$

Sequential Bayes': $p(x \mid z_{1:t}) = \frac{\eta_t}{\eta_{t-1}} p(z_t \mid x)p(x \mid z_{1:t-1})$

Occupancy Grid Bayesian Update for Cell (i,j)

$$\bar{p}(m(i,j) | z_{1:t}(i,j)) = p(z_t(i,j) | m(i,j))p(m(i,j) | z_{1:t-1}(i,j))$$

Occ Prob at t (unnormalized) Occ Prob at t-1

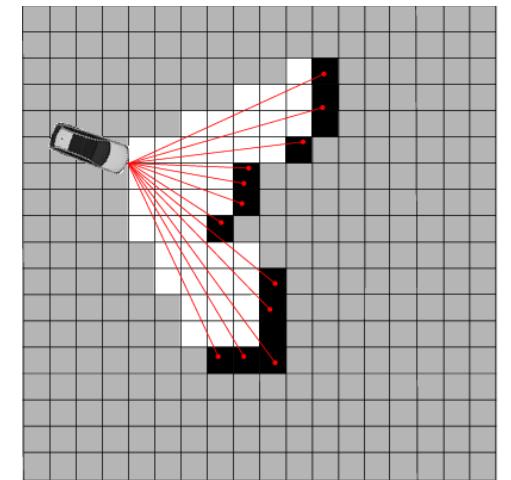
See point in cell Cell occupied

$$p(z_t(i,j) | m(i,j)) = \begin{bmatrix} \text{True negative} & \text{False negative} \\ p(0 | 0) & p(0 | 1) \\ p(1 | 0) & p(1 | 1) \end{bmatrix}$$

False positive True positive

Occ Prob at t (normalized)

$$p(m(i,j) | z_{1:t}(i,j)) = \frac{\bar{p}(m(i,j) | z_{1:t-1}(i,j))}{\bar{p}(m(i,j) = 0 | z_{1:t-1}(i,j)) + \bar{p}(m(i,j) = 1 | z_{1:t-1}(i,j))}$$

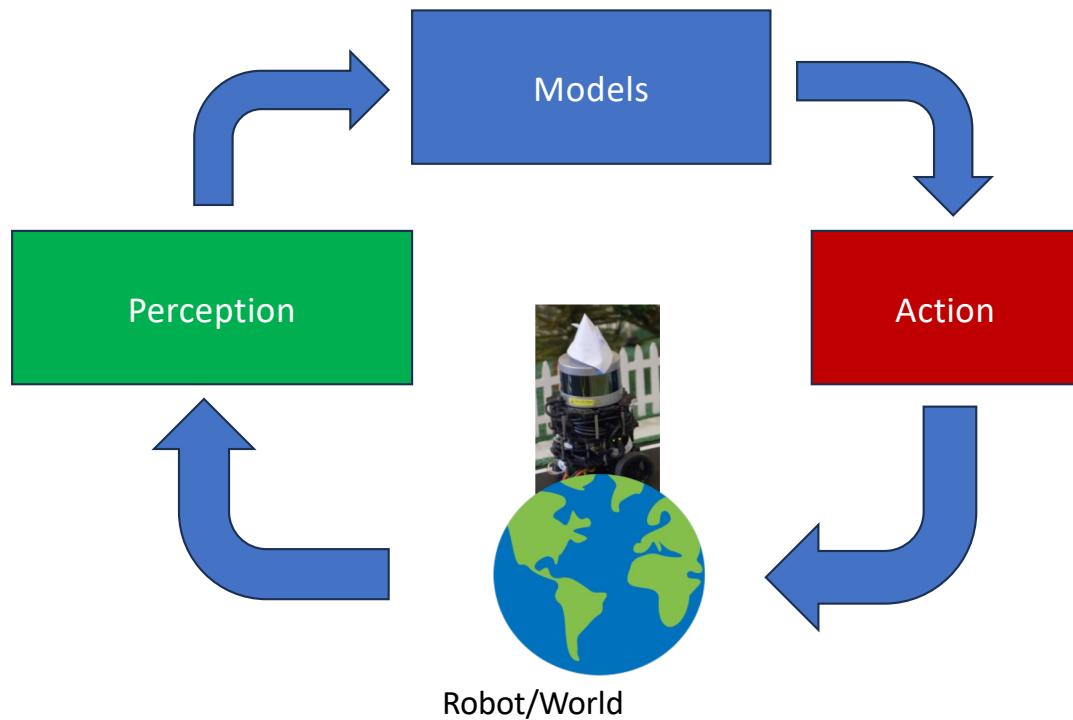


Expectation of a RV

- Expectation for discrete RVs: $E[X] = \sum_x x p(x)$
- Expectation for continuous RVs: $E[X] = \int x p(x) dx$
- Expectation is a linear operator: $E[aX + b] = a E[X] + b$
- Expectation of a vector of RVs is simply the vector of expectations
- Covariance

$$\text{cov}(X, Y) = E[(X - E[X])(Y - E[Y])^T] = E[XY^T] - E[X]E[Y]^T$$

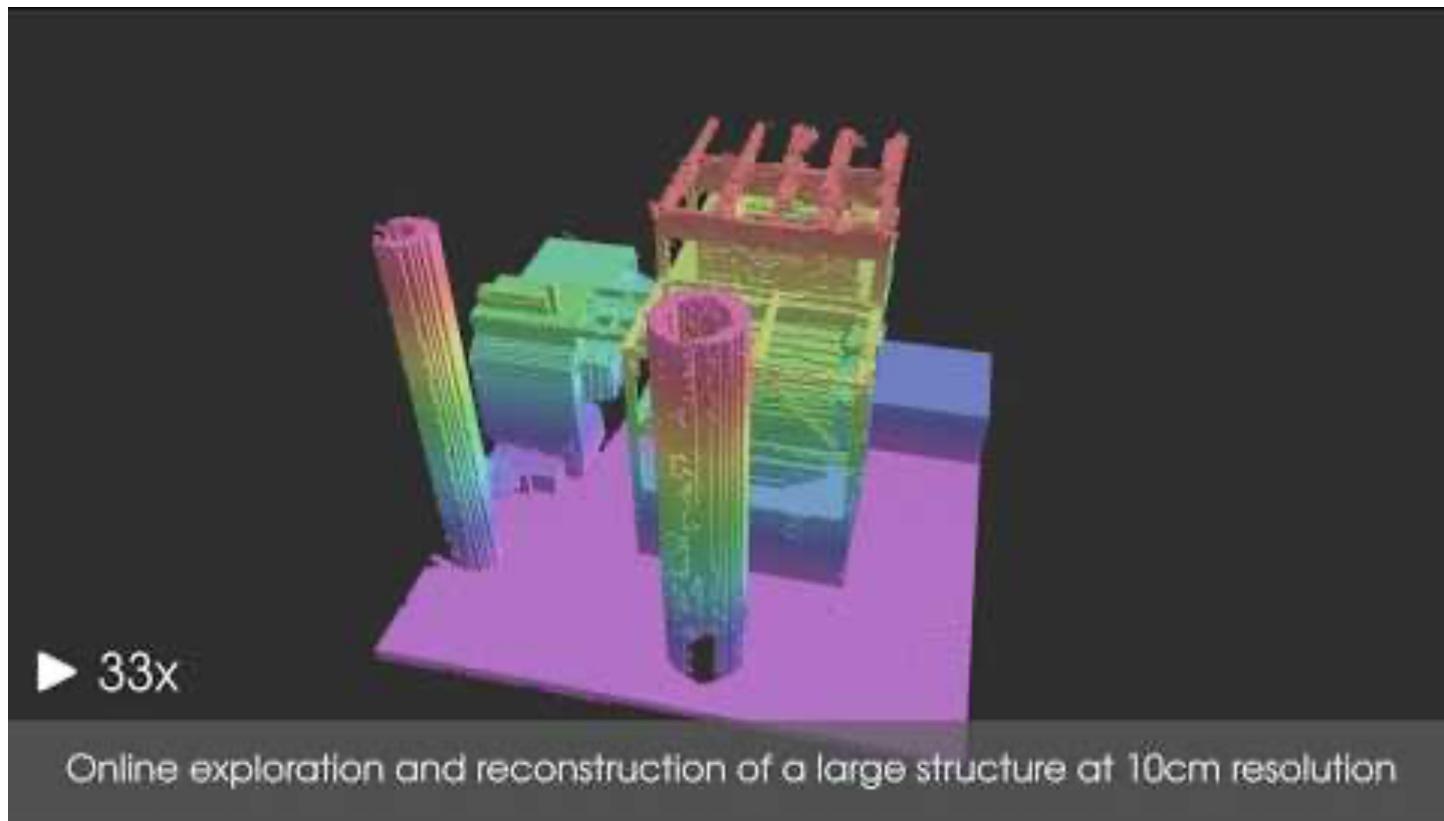
Active Perception as a Planning Objective



Active Perception:
Take actions to improve perception performance

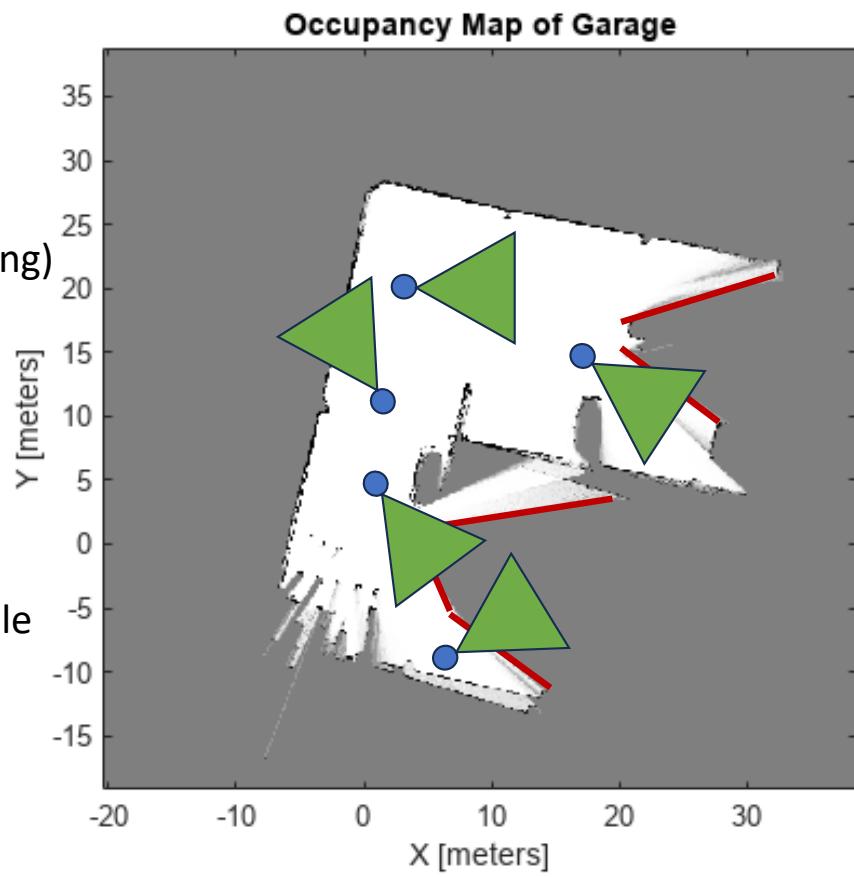
- Exploration (to improve a map)
- Informative planning (to improve robot pose estimate)
- Target following (to improve target pose estimate)

Frontier Exploration



Frontier Exploration

- Frontier points:
 - High uncertainty (prob occ ~ 0.5)
 - Neighboring other cells with low uncertainty
- Consider robot pose candidates (e.g. from sampling)
- Score poses based on:
 - Number of frontier cells in footprint
 - Length of path to get there
- Pick the best and execute
- Called next-best-view planning, myopic planning, Greedy planning
- Other methods solve an optimization over multiple future view points
- “Orienteering problem”



Information gathering, informative planning, active sensing, belief space planning

- Probabilistic approach
- Quantify “information gain” from a Bayesian estimate
- Can be
 - A map
 - Target state
 - Ego state of the robot
- Turn into trajectory optimization problem: find sequence of control inputs to maximize expected information gain over a time horizon
- Main advantage: general unifying approach
- Main drawback: tends to be computationally expensive

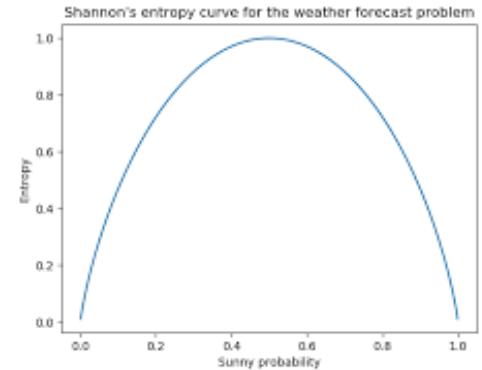
Information gathering, informative planning, active sensing, belief space planning

E.g., binary RV:

Shannon Entropy: measure of uncertainty in a RV

Discrete: $H(X) = - \sum_x \log_2(p(x))p(x)$

Continuous: $H(X) = - \int_x \ln(p(x))p(x)dx$



E.g., Gaussian RV:

$$H(X) = C + \frac{1}{2} \ln \det(\Sigma)$$

Information gathering, informative planning, active sensing, belief space planning

Conditional Entropy: Expected entropy after taking a measurement

$$\begin{aligned}
 H(X | Y) &= \mathbb{E}[H(X) | Y] = \int_y \left(- \int_x p(x | y) \ln(p(x | y)) dx \right) p(y) dy \\
 &= - \int_y \int_x p(y | x)p(x) \ln\left(\frac{p(y | x)p(x)}{p(y)}\right) dxdy
 \end{aligned}$$

Measurement likelihood
 Depends on robot state
 Prior

Mutual information: Expected entropy reduction
after taking a measurement

$$I(X, Y) = H(X) - H(X | Y)$$

Informative planning: Plan trajectory for robot
to maximize mutual information

Challenging to solve in practice!

$$\max_{u_{1:T}} \sum_{t=1}^T I(X, Y; s_t) - u_t^T R u_t$$

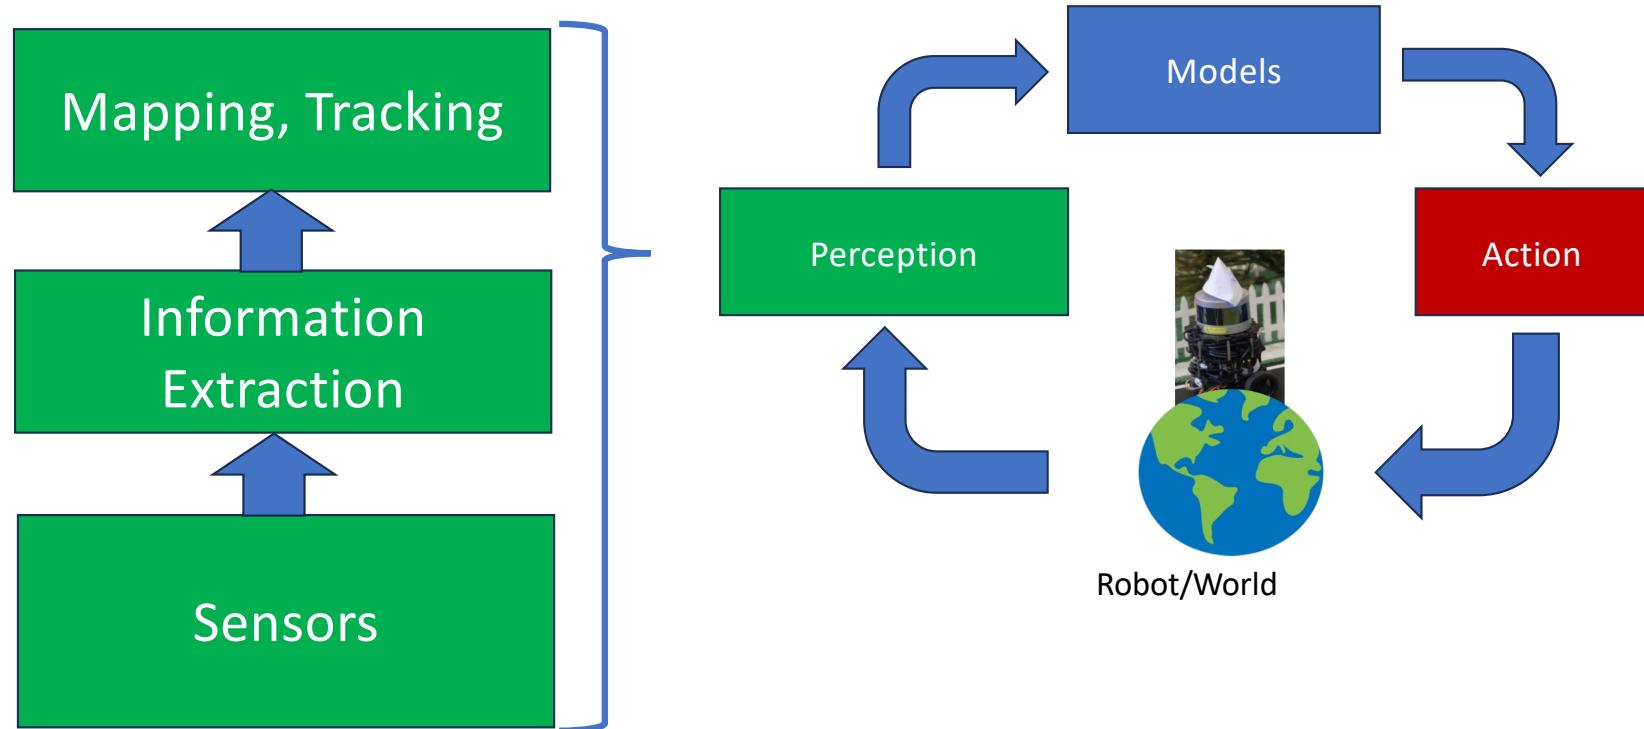
Dynamics $\xrightarrow{\text{s.t.}} s_{t+1} = f(s_t, u_t)$
 Collision constraints $\xrightarrow{} g(s_t) \leq 0$
 Input constraints $\xrightarrow{} h(u_t) \leq 0$

Information gathering, informative planning, active sensing, belief space planning

- Highly versatile and general
- Works for any problem posed as a Bayesian filter
 - Special variants for occ maps, EKFs, particle filters, UKFs
 - Special variants for map exploration, robot ego state estimation, active target tracking and target search, active SLAM
- Drawback: usually requires approximations and computational shortcuts

Robot Perception

Perception Stack



Model for robot-environment interaction

- Two fundamental types of robot-environment interactions: the robot can influence **the state** of its environment through **control actions**, and gather information about the **state** through **measurements**
- **State x_t** : collection at time t of all aspects of the robot and its environment that can impact the future
 - Map (location of fixed objects in environment)
 - Robot pose (e.g., robot location and orientation)
 - Robot velocity
 - State of dynamics objects in environment
- Useful notation: $x_{t_1:t_2} := x_{t_1}, x_{t_1+1}, x_{t_1+2}, \dots, x_{t_2}$
- A state x_t is called *complete* if no variables prior to x_t can influence the evolution of future states → **Markov property**

Measurement and control data

- **Measurement data z_t :** information about state of the environment at time t ; useful notation

$$z_{t_1:t_2} := z_{t_1}, z_{t_1+1}, z_{t_1+2}, \dots, z_{t_2}$$

- **Control data u_t :** information about the change of state at time t ; useful notation

$$u_{t_1:t_2} := u_{t_1}, u_{t_1+1}, u_{t_1+2}, \dots, u_{t_2}$$

- Key difference: measurement data tends to increase robot's knowledge, while control actions tend to induce a loss of knowledge

State equation

- General probabilistic generative model

$$p(x_t | x_{0:t-1}, z_{1:t-1}, u_{1:t})$$

Convention: first take control action and then take measurement

- **Key assumption:** state is complete (i.e., the Markov property holds)

State transition probability $\longrightarrow p(x_t | x_{0:t-1}, z_{1:t-1}, u_{1:t}) = p(x_t | x_{t-1}, u_t)$

- In other words, we assume *conditional independence*, with respect to conditioning on x_{t-1}
- Special case (typical dynamics model):

$$x_t = f(x_{t-1}, u_{t-1}) + w_{t-1}, \quad w_{t-1} \sim \mathcal{N}(0, Q_{t-1})$$

Measurement equation and overall stochastic model

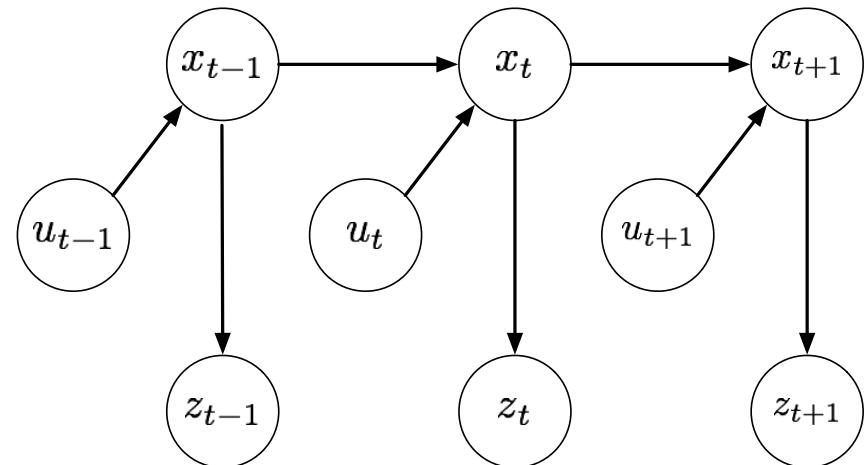
- Assuming x_t is complete

$$\xrightarrow{\text{Measurement probability}} p(z_t | x_{0:t}, z_{1:t-1}, u_{1:t}) = p(z_t | x_t)$$

Measurement probability

- Overall dynamic Bayes network model (also referred to as hidden Markov model)
- Special case (typical measurement model):

$$z_t = g(x_t, u_t) + v_t, \quad v_t \sim \mathcal{N}(0, R_t)$$



Belief distribution

- **Belief distribution:** reflects internal knowledge about the state
- A belief distribution assigns a probability to each possible hypothesis about the true state
- Formally, belief distributions are posterior probabilities over state variables conditioned on the available data

$$bel(x_t) := p(x_t | z_{1:t}, u_{1:t})$$

- Similarly, the *prediction* distribution is defined as

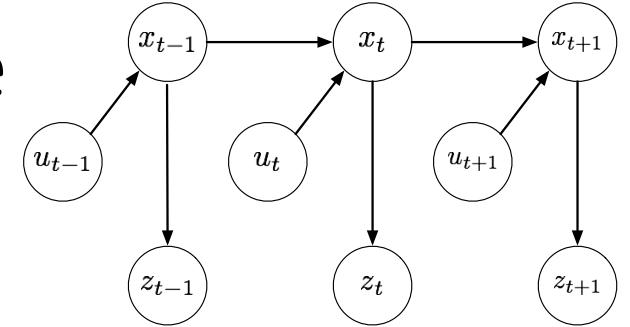
$$\overline{bel}(x_t) := p(x_t | \mathbf{\tilde{z}}_{1:t-1}, u_{1:t})$$

- Calculating $bel(x_t)$ from $\overline{bel}(x_t)$ is called correction or measurement update

Bayes filter algorithm

- Bayes' filter algorithm: most general algorithm for calculating beliefs
 - Key assumption: state is complete
 - Recursive algorithm
 - Step 1 (prediction): compute $\bar{bel}(x_t)$
 - Step 2 (measurement update): compute $bel(x_t)$
 - Algorithm initialized with $bel(x_0)$ (e.g., uniform or points mass)
- Data:** $bel(x_{t-1}), u_t, z_t$
Result: $bel(x_t)$
foreach x_t **do**
- $$\left| \begin{array}{l} \bar{bel}(x_t) = \int p(x_t | u_t, x_{t-1}) bel(x_{t-1}) dx_{t-1}; \\ bel(x_t) = \eta p(z_t | x_t) \bar{bel}(x_t); \end{array} \right.$$
- end**
Return $bel(x_t)$

Derivation: measurement update



$$bel(x_t) = p(x_t \mid z_{1:t}, u_{1:t})$$

$$= \frac{p(z_t \mid x_t, z_{1:t-1}, u_{1:t}) p(x_t \mid z_{1:t-1}, u_{1:t})}{\underbrace{p(z_t \mid z_{1:t-1}, u_{1:t})}_{:=\eta^{-1}}} \quad \text{Bayes rule}$$

$$= \eta p(z_t \mid x_t) \underbrace{p(x_t \mid z_{1:t-1}, u_{1:t})}_{= \overline{bel(x_t)}} \quad \text{Markov property}$$

Derivation: prediction step

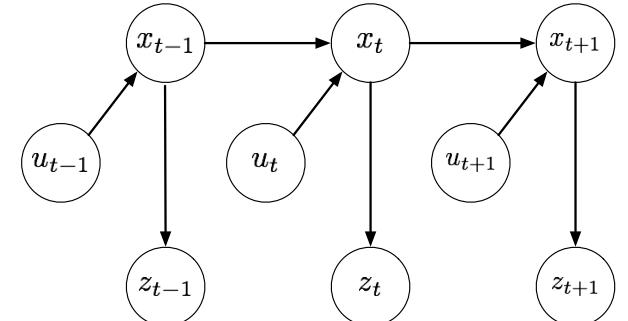
$$\overline{bel}(x_t) = p(x_t | z_{1:t-1}, u_{1:t})$$

$$= \int p(x_t | x_{t-1}, z_{1:t-1}, u_{1:t}) p(x_{t-1} | z_{1:t-1}, u_{1:t}) dx_{t-1}$$

$$= \int p(x_t | x_{t-1}, u_t) p(x_{t-1} | z_{1:t-1}, u_{1:t}) dx_{t-1}$$

$$= \int p(x_t | x_{t-1}, u_t) p(x_{t-1} | z_{1:t-1}, \textcolor{red}{u_{1:t-1}}) dx_{t-1}$$

$$= \int p(x_t | x_{t-1}, u_t) bel(x_{t-1}) dx_{t-1}$$



Total probability

Markov

For general output feedback policies, u_t does not provide additional information on x_{t-1}

Discrete Bayes' filter

- Discrete Bayes' filter algorithm: applies to problems with *finite* state spaces

- Belief $bel(x_t)$ represented as pmf $\{p_{k,t}\}$

Data: $\{p_{k,t-1}\}, u_t, z_t$
Result: $\{p_{k,t}\}$
foreach k **do**
 $\bar{p}_{k,t} = \sum_i p(X_t = x_k | u_t, X_{t-1} = x_i) p_{i,t-1};$
 $p_{k,t} = \eta p(z_t | X_t = x_k) \bar{p}_{k,t};$
end
Return $\{p_{k,t}\}$

Simple example – Belief & Measurement Model

Adapted from [PR]

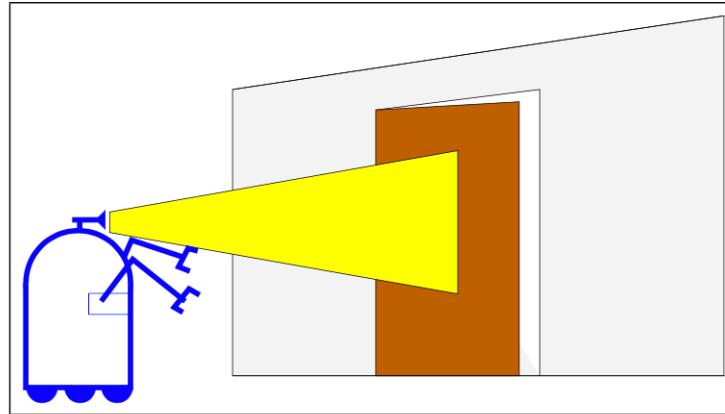


Figure 2.2 A mobile robot estimating the state of a door.

$$bel(X_0 = \text{open}) = 0.5$$

$$bel(X_0 = \text{closed}) = 0.5$$

$$p(Z_t = \text{sense_open} | X_t = \text{is_open}) = 0.6$$

$$p(Z_t = \text{sense_closed} | X_t = \text{is_open}) = 0.4$$

$$p(Z_t = \text{sense_open} | X_t = \text{is_closed}) = 0.2$$

$$p(Z_t = \text{sense_closed} | X_t = \text{is_closed}) = 0.8$$

Simple example – Transition Model

Adapted from [PR]

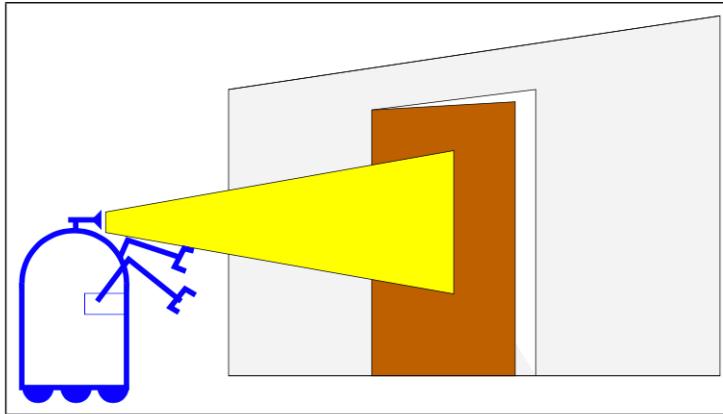


Figure 2.2 A mobile robot estimating the state of a door.

$$\begin{array}{lll} p(X_t = \text{is_open} | U_t = \text{push}, X_{t-1} = \text{is_open}) & = & 1 \\ p(X_t = \text{is_closed} | U_t = \text{push}, X_{t-1} = \text{is_open}) & = & 0 \\ p(X_t = \text{is_open} | U_t = \text{push}, X_{t-1} = \text{is_closed}) & = & 0.8 \\ p(X_t = \text{is_closed} | U_t = \text{push}, X_{t-1} = \text{is_closed}) & = & 0.2 \end{array} \quad \begin{array}{lll} p(X_t = \text{is_open} | U_t = \text{do_nothing}, X_{t-1} = \text{is_open}) & = & 1 \\ p(X_t = \text{is_closed} | U_t = \text{do_nothing}, X_{t-1} = \text{is_open}) & = & 0 \\ p(X_t = \text{is_open} | U_t = \text{do_nothing}, X_{t-1} = \text{is_closed}) & = & 0 \\ p(X_t = \text{is_closed} | U_t = \text{do_nothing}, X_{t-1} = \text{is_closed}) & = & 1 \end{array}$$

Simple Example – Update Step

Adapted from [PR]

- Is door open or not?

$$bel(X_0 = \text{open}) = 0.5$$

$$bel(X_0 = \text{closed}) = 0.5$$

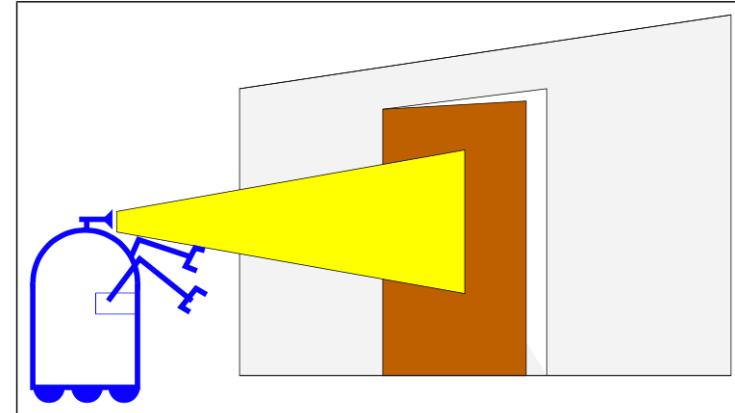


Figure 2.2 A mobile robot estimating the state of a door.

Measurement Model

$$\begin{aligned} p(Z_t = \text{sense_open} | X_t = \text{is_open}) &= 0.6 \\ p(Z_t = \text{sense_closed} | X_t = \text{is_open}) &= 0.4 \end{aligned}$$

$$\begin{aligned} p(Z_t = \text{sense_open} | X_t = \text{is_closed}) &= 0.2 \\ p(Z_t = \text{sense_closed} | X_t = \text{is_closed}) &= 0.8 \end{aligned}$$

Transition Model for do_nothing

$$\begin{aligned} p(X_t = \text{is_open} | U_t = \text{do_nothing}, X_{t-1} = \text{is_open}) &= 1 \\ p(X_t = \text{is_closed} | U_t = \text{do_nothing}, X_{t-1} = \text{is_open}) &= 0 \\ p(X_t = \text{is_open} | U_t = \text{do_nothing}, X_{t-1} = \text{is_closed}) &= 0 \\ p(X_t = \text{is_closed} | U_t = \text{do_nothing}, X_{t-1} = \text{is_closed}) &= 1 \end{aligned}$$

Simple Example – Update Step

Adapted from [PR]

- Is door open or not?

$$bel(X_0 = \text{open}) = 0.5$$

$$bel(X_0 = \text{closed}) = 0.5$$

Received Sensor Measurement:

$$bel(x_1) = \eta p(Z_1 = \text{sense_open} | x_1) \overline{bel}(x_1)$$

Measurement Model

$$p(Z_t = \text{sense_open} | X_t = \text{is_open}) = 0.6$$

$$p(Z_t = \text{sense_closed} | X_t = \text{is_open}) = 0.4$$

$$p(Z_t = \text{sense_open} | X_t = \text{is_closed}) = 0.2$$

$$p(Z_t = \text{sense_closed} | X_t = \text{is_closed}) = 0.8$$

$bel(x_t) = (0.75, 0.25)$ Don't forget normalization!

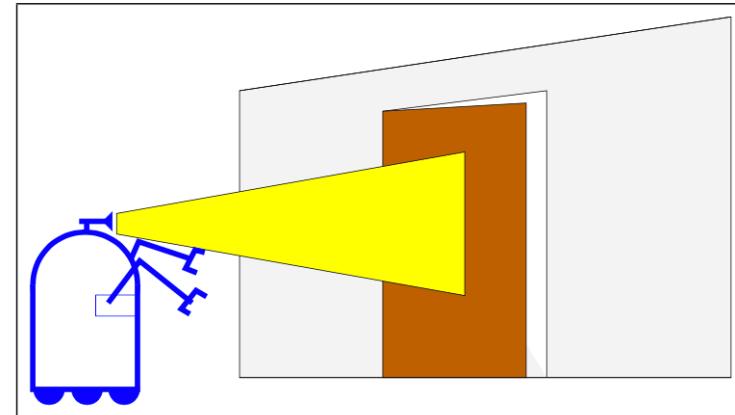


Figure 2.2 A mobile robot estimating the state of a door.

Transition Model for do_nothing

$$p(X_t = \text{is_open} | U_t = \text{do_nothing}, X_{t-1} = \text{is_open}) = 1$$

$$p(X_t = \text{is_closed} | U_t = \text{do_nothing}, X_{t-1} = \text{is_open}) = 0$$

$$p(X_t = \text{is_open} | U_t = \text{do_nothing}, X_{t-1} = \text{is_closed}) = 0$$

$$p(X_t = \text{is_closed} | U_t = \text{do_nothing}, X_{t-1} = \text{is_closed}) = 1$$

Next time

